

MongoDB and CouchDB

Instructor: Dr. GP Saggese - gsaggese@umd.edu

- References:
 - All concepts in slides
 - MongoDB tutorial
 - Web
 - <https://www.mongodb.com/>
 - Official docs
 - pymongo
- Book
 - Seven Databases in Seven Weeks, 2e




Key-Value Store vs Document DBs

- **Key-value stores**

- Function as a map or dictionary
 - Examples: HBase, Redis
- Primarily retrieve values using keys
 - Occasionally search within value fields using patterns
- Store uninterpreted values (e.g., binary blobs) linked to keys
- Use a single namespace for all key-value pairs


Key-Value



- **Document DBs**

- Group key-value pairs into *documents*
 - Examples: MongoDB, CouchDB
- Documents formatted in JSON, XML, or BSON (binary JSON)
- Documents are part of *collections*
 - Comparable to *tables* in relational databases
 - Large collections can be partitioned and indexed

Document



MongoDB



- Developed by MongoDB Inc.
 - Founded in 2007
 - Based on DoubleClick experience with large-scale data
 - Mongo comes from “hu-mongo-us”
- Highly popular NoSQL database
- **Document-oriented NoSQL DB**
 - Schema-less
 - No Data Definition Language (DDL) like SQL
 - Store maps with any keys and values
 - Application manages schema, linking documents to meanings
 - Keys are string-stored hashes
 - Each document has a unique `_id` (reserved by Mongo)
 - Values in BSON format
 - Based on JSON (B for Binary)
- Developed in C++
- Supports APIs (drivers) in various languages
 - Examples: JavaScript, Python, Ruby, Java, Scala, C++, etc

MongoDB: Example of Document

- A document is a JSON data structure
- It corresponds to a row in a relational DB
 - Without schema
 - Primary key is `_id`
 - Values nested to an arbitrary depth

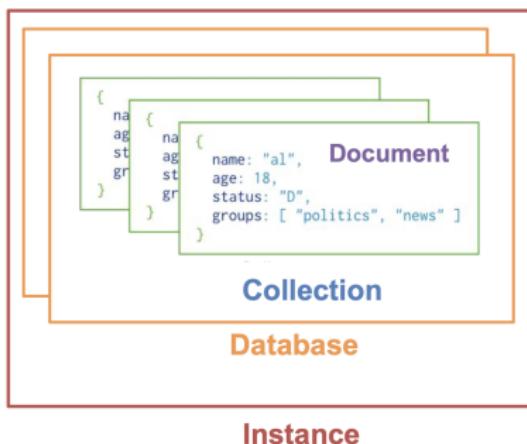
```
{  
    "_id" : ObjectId("4d0b6da3bb30773266f39fea"),  
    "country" : {  
        "$ref" : "countries",  
        "$id" : ObjectId("4d0e6074deb8995216a8309e")  
    },  
    "famous_for" : [  
        "beer",  
        "food"  
    ],  
    "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",  
    "mayor" : {  
        "name" : "Ted Wheeler",  
        "party" : "D"  
    },  
    "name" : "Portland",  
    "population" : 582000,  
    "state" : "OR"  
}
```

MongoDB: Functionalities

- **Design goals**
 - Performance
 - Availability/scalability
 - Rich data storage (not rich querying!)
- **Dynamic schema**
 - No DDL
 - Secondary indexes
 - Query language via API
- **Several levels of data consistency**
 - Atomic writes and fully-consistent reads (document level)
- **No joins nor transactions across multiple documents**
 - Distributed queries easy and fast
- **High availability through replica sets**
 - Primary replication with automated failover
- **Built-in sharding**
 - Horizontal scaling via automated range-based partitioning
 - Reads and writes distributed over shards

MongoDB: Hierarchical Objects

- A Mongo **instance** has:
 - Zero or more “databases”
 - Mongo instance ~ Postgres instance
- A Mongo **database** has:
 - Zero or more “collections”
 - Mongo collection ~ Postgres tables
 - Mongo database ~ Postgres database
- A Mongo **collection** has:
 - Zero or more “documents”
 - Mongo document ~ Postgres rows
- A Mongo **document** has:
 - One or more “fields”
 - Always has primary key `_id`
 - Mongo field ~ Postgres columns



Relational DBs vs MongoDB: Concepts

RDBMS Concept	MongoDB Concept	Meaning in MongoDB
database	database	Container for collections
relation / table / view	collection	Group of documents
row / instance	document	Group of fields
column / attribute	field	A name-value pair
index	index	Automatic
primary keys	_id field	Always the primary key
foreign key	reference	Pointers
table joins	embedded documents	Nested name-value pairs

```
{  
    "_id" : ObjectId("4d0b6da3bb30773266f39fea"),  
    "country" : {  
        "$ref" : "countries",  
        "$id" : ObjectId("4d0e6074deb8995216a8309e")  
    },  
    "famous_for" : [  
        "beer",  
        "food"  
    ],  
    "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",  
    "mayor" : {  
        "name" : "Ted Wheeler",  
        "party" : "D"  
    },  
    "name" : "Portland",  
    "population" : 582000,  
    "state" : "OR"  
}
```

Relational vs Document DB: Workflows

- **Relational DBs**


- E.g., PostgreSQL
- Know what to store
 - Tabular data
- Static schema allows query flexibility (e.g., joins)
- Complexity at insertion time
 - Decide data representation (schema)

- **Document DBs**

- E.g., MongoDB
- No assumptions on storage
 - E.g., irregular JSON data
- Access data by key
 - Nested key-value map
- Complexity at access time
 - Retrieve data from server
 - Process data client-side

Why Use MongoDB?

- Simple to query
 - Work on client side
- Fast
 - 2-10x faster than Postgres
- Data model/functionalities suitable for most web applications
 - Semi-structured data
 - Quickly evolving systems
- Easy and fast data integration
- Not suited for heavy, complex transaction systems
 - E.g., banking system



MongoDB: Data Model


- **Documents** are field-value pairs
 - **Field names**: strings
 - **Values**: any BSON type
 - Arrays of documents
 - Native data types
 - Other documents
- Examples:
 - `_id: ObjectId`
 - `name: document with fields first and last`
 - `birth and death: date type`
 - `contribs: array of strings`
 - `views: NumberLong type`

```
{  
    name: "sue",           ← field:value  
    age: 26,               ← field:value  
    status: "A",           ← field:value  
    groups: [ "news", "sports" ] ← field:value  
}  
  
{  
    _id: ObjectId("5099803df3f4948bd2f98391"),  
    name: { first: "Alan", last: "Turing" },  
    birth: new Date('Jun 23, 1912'),  
    death: new Date('Jun 07, 1954'),  
    contribs: [ "Turing machine", "Turing test", "Turingery" ],  
    views : NumberLong(1250000)  
}
```

MongoDB: Data Model

- Documents can be nested
 - Embedded sub-document
- **Denormalized data models**
 - Store related information in the same record
 - Result of a join operation
- **Normalized data models**
 - Eliminate duplication
 - Represent many-to-many relationships

```
{  
  _id: <ObjectId1>,  
  username: "123xyz",  
  contact: {  
    phone: "123-456-7890",  
    email: "xyz@example.com"  
  },  
  access: {  
    level: 5,  
    group: "dev"  
  }  
}
```



user document

```
{  
  _id: <ObjectId1>,  
  username: "123xyz"  
}
```

contact document

```
{  
  _id: <ObjectId2>,  
  user_id: <ObjectId1>,  
  phone: "123-456-7890",  
  email: "xyz@example.com"  
}
```

access document

```
{  
  _id: <ObjectId3>,  
  user_id: <ObjectId1>,  
  level: 5,  
  group: "dev"  
}
```

Schema Free

- MongoDB does not need pre-defined data schema
- Every **document** in a **collection** can have different fields and values
 - No need for NULL values / union of fields like in relational DBs
- E.g., dishomogeneous data instances

```
{name: "will",
  eyes: "blue",
  birthplace: "NY",
  aliases: ["bill", "la ciacco"],
  loc: [32.7, 63.4],
  boss: "ben"}
```

```
{name: "jeff",
  eyes: "blue",
  loc: [40.7, 73.4],
  boss: "ben"}
```

```
{name: "brendan",
  aliases: ["el diablo"]}
```

```
{name: "ben",
  hat: "yes"}
```

```
{name: "matt",
  pizza: "DiGiorno",
  height: 72,
  loc: [44.6, 71.3]}
```



mongoDB



SCIENCE
ACADEMY

JSON Format

- JSON = JavaScript Object Notation

- Data is stored in field / value pairs
- A field / value pair consists of:
 - A field name (always a string)
 - Followed by a colon :
 - Followed by a typed value

`"name": "R2-D2"`

- Data in documents is separated by commas ','

`"name": "R2-D2", race: "Droid"`

- Curly braces {} hold documents

`{"name": "R2-D2", race : "Droid", affiliation: "rebels"}`

- An array is stored in brackets []

`[{"name": "R2-D2", race: "Droid", affiliation: "rebels"}, {"name": "Yoda", affiliation: "rebels"}]`

- Supports:

SCIENCE Embedding of nested objects within other objects
ACADEMY Just references




BSON Format

- Binary-encoded serialization of JSON-like documents
 - <https://bsonspec.org>
- Store zero or more key/value pairs as a single entity
 - Each entry consists of:
 - Field name (string)
 - Data type
 - Value
- Similar to Protocol Buffer, but more schema-less
- Prefix large elements in BSON document with a length field for scanning
- MongoDB understands BSON objects, even nested
 - Build indexes and match objects against query expressions for BSON keys

ObjectId

- Each JSON data contains an `_id` field of type `ObjectId`
 - Similar to `SERIAL` constraint incrementing a numeric primary key in PostgreSQL
- An `ObjectId` is 12 bytes, composed of:
 - Timestamp
 - Client machine ID
 - Client process ID
 - 3-byte auto-incremented counter
- Each Mongo process handles its own ID generation without colliding
 - Mongo's distributed nature
- Details [here](#)



Indexes

- Primary index
 - Automatically created on _id field
 - B+ tree indexes
- Secondary index
 - Improve query performance
 - Enforce unique values for a field
- Single field index and compound index (like SQL)
 - Order of fields in a compound index matters
- Sparse property of an index
 - Index contains entries for documents with the indexed field
 - Ignore records without the field
- Reject records with duplicate key if index is unique and sparse
- Details [here](#)