

8.2: Map Reduce

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu
- **References**
 - Silberschatz: Chap 10
 - Ghemawat et al.: *The Google File System*, 2003
 - Dean et al.: *MapReduce: Simplified Data Processing on Large Clusters*, 2004

SEVENTH EDITION
Database System Concepts



MapReduce: Overview

- **MapReduce programming model**

- Inspired by functional programming (e.g., Lisp)
- Common pattern of parallel programming to process large number of records

- **Basic algorithm**

- Apply `map()` to each record
- Group results by key
- Apply `reduce()` to results of `map()`

- **Example**

- *Goal:* Sum length of all tuples in a document
 - E.g.,
[(), (a,), (a, b), (a, b, c)]
- *map(function, set of values)*
 - Apply function to each value (e.g., `len`)
`map(len, [(), (a,), (a, b), (a, b, c)])` -> [0, 1, 2, 3]
- *reduce(function, set of values)*
 - Combine values using a binary function (e.g., `add`)
`reduce(add, [0, 1, 2, 3])` -> 6

MapReduce: Overview

- **Structure of computation**

- *Read input*
 - Sequentially or in parallel
- *Map*
 - Extract / compute from records
- *Group by key*
 - Sort and shuffle
- *Reduce*
 - Aggregate, summarize, filter, transform
- *Write result*

- **Division of responsibilities**

- User specifies `map()` and `reduce()` functions to solve problem
- MapReduce framework (e.g., Hadoop, Spark) implements algorithm

MapReduce: Word Count

- **Word Count**

- “Hello world” of MapReduce
- Huge text file (can’t fit in memory)
- Count occurrences of each distinct word

- **Linux solution**

```
> more doc.txt
```

One a penny, two a penny, hot cross buns.

```
> words doc.txt | sort | uniq -c
```

a 2

buns 1

cross 1

...

- words outputs words one per line
- Unix pipeline is parallelizable in MapReduce sense

Hot cross buns!

Hot cross buns!

One a penny, two a penny,

Hot cross buns!

If you have no daughters,
Give them to your sons.

One a penny, two a penny,
Hot cross buns!^[1]

- **Sample application**

- Analyze web server logs for popular URLs

MapReduce: Word Count

Action

Read input

Map:

- Invoke `map()` on each input record
- Emit 0 or more output data items

Group by key:

- Gather all outputs from `map()` stage
- Collect outputs by keys

Reduce:

- Combine the list of outputs with same keys

Python code

```
values = read(file_name)
```

```
def map(values):
```

values: words in document

```
for word in values:
```

```
    emit(word, 1)
```

```
def reduce(key, values):
```

key: a word

value: a list of counts

```
result = 0
```

result = sum(values)

```
for count in values:
```

```
    result += count
```

```
emit(key, result)
```

Example

"One a penny, two a penny,
hot cross buns."

Map:

```
[("one", 1), ("a", 1),  
("penny", 1), ("two", 1),  
("a", 1), ("penny", 1),  
("hot", 1), ("cross", 1),  
("buns", 1)]
```

Group by key:

```
[("a", [1, 1]),  
("buns", [1]),  
("cross", [1]),  
("hot", [1]),  
("one", [1]),  
("penny", [1, 1]),  
("two", [1])]
```

Reduce:

```
[("one", 1),  
("a", 2),  
("penny", 2),  
("two", 1),  
("hot", 1),  
("cross", 1),  
("buns", 1)]
```

MapReduce: Log Processing

- **Goal:
 - Log file recording access to a website with format (date, hour, filename)
 - Find how many times each file is accessed during Feb 2013
 - **Input**
 - Read file and split into lines
 - **Map**
 - Parse each line into 3 fields
 - If date is in the required interval
`emit(dir_name, 1)`
 - **GroupBy**
 - Reduce key is the filename
 - Accumulate all (key, value) with the same filename
 - **Reduce**
 - Add values for each list of (key, value) with the same filename
 - Output number of accesses to each file
 - **Output**
 - Write results on disk separated by ...
- After Input*
- ```
2013/02/21 10:31:22.00EST /slide-
2013/02/21 10:43:12.00EST /slide-
2013/02/22 18:26:45.00EST /slide-
2013/02/22 18:26:48.00EST /exer-
2013/02/22 18:26:54.00EST /exer-
2013/02/22 20:53:29.00EST /slide-
```
- After Map*
- ```
['/slide-dir/11.ppt', 1], ...])
```
- After GroupBy*
- ```
[('/slide_dir/11.ppt', 1), ..., ('slide-dir/12.ppt', [1, 1]), ...]
```
- After Reduce*
- ```
[('/slide_dir/11.ppt', 1), ..., ('slide-dir/12.ppt', 2), ...]
```
- Output*
- ```
/slide_dir/11.ppt 1
...
/slide-dir/12.ppt 2
...
```



# MapReduce: Interfaces

---

- **Input:**

- read key-value pairs **List[Tuple[k, v]]**
- Programmer specifies two methods `map` and `reduce`
- **Map(Tuple[k, v]) → List[Tuple[k, v]]**

- Take a key-value pair and output a set of key-value pairs
  - E.g., key is a file, value is the number of occurrences
  - “One a penny” → [ (“One”, 1), (“a”, 1), (“penny”, 1) ]
- There is one **Map** call for every  $(k, v)$  pair

- **GroupBy(List[Tuple[k, v]]) → List[Tuple[k, List[v]]]**

- Group and optionally sort all the records with the reduce key

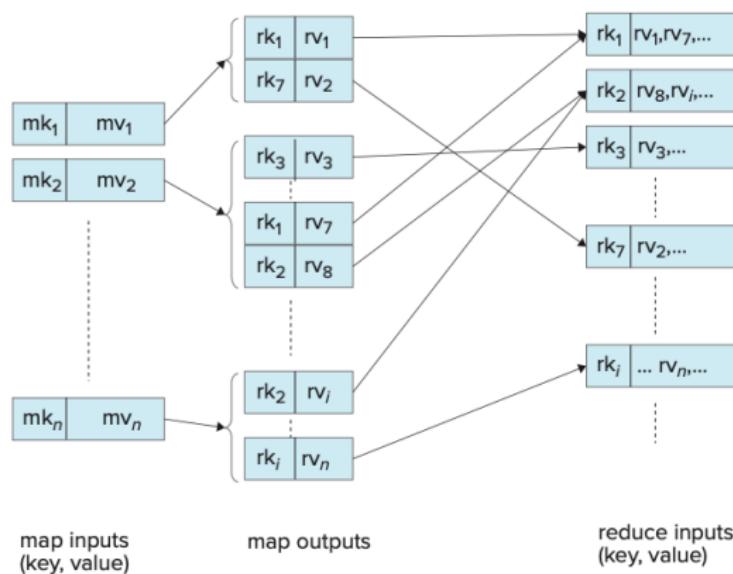
- **Reduce(Tuple[k, List[v]]) → Tuple[k, v]**

- All values  $v'$  with same key  $k'$  are reduced together
- There is one **Reduce** call per unique key  $k'$

- Output: write key-value pairs **List[Tuple[k, v]]**

# MapReduce: Data Flow

Focusing on MapReduce functionality / flow of the data to expose the parallelism



Input      Map

GroupBy    Reduce .

- **Input**
- **Map**

- $mki$  = map keys
- $mvi$  = map input values

- **GroupBy**

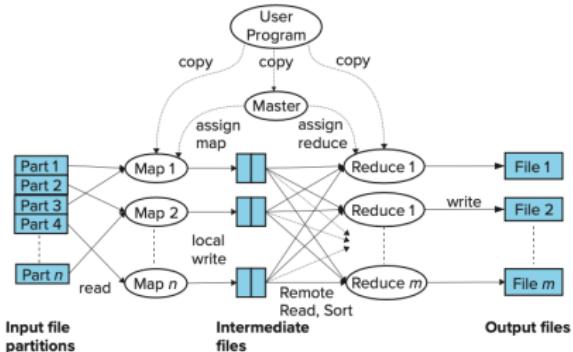
- Shuffle / collect the data

- **Reduce**

- $rki$  = reduce keys
- $rvi^{**}$  = reduce input values
- Reduce outputs are not shown

# MapReduce: Parallel Data Flow

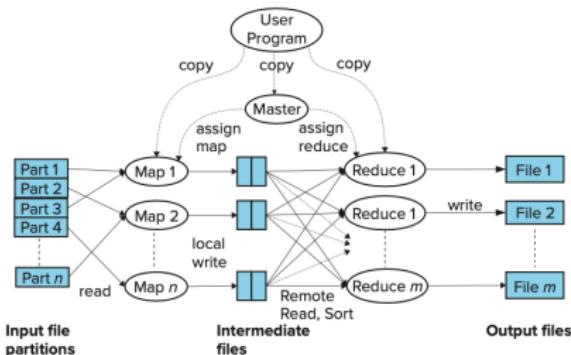
- **User program** specifies map/reduce code
- **Input data** is partitioned across multiple machines (HDFS)
- **Master** node sends copies of the code to all computing nodes
- **Map**
  - $n$  data chunks to process
  - Functions executed in parallel on multiple  $k$  machines
  - Output data from *Map* is saved on disk
- **GroupBy / Sort**
  - Output data from *Map* is sorted and partitioned based on reduce key
  - Different files are created for each *Reduce* task
- **Reduce**
  - Functions executed in parallel on multiple machines
  - Each work on some part of the data
  - Output data from *Reduce* is saved on disk



- All operations use HDFS as storage
- Machines are reused for multiple computations (Map, GroupBy, Reduce) at different times

# Master Node Responsibilities

- **Master node coordinates**
  - Task status: idle, in-progress, completed
  - Schedule idle tasks as workers become available
  - Map task completion sends location and sizes of intermediate files to Master
  - Master informs Reduce tasks
  - Schedule idle Reduce tasks
- **Master node pings workers to detect failures**



# Dealing with Failures

---

- **Map worker failure**
  - Reset failed map tasks to idle
  - Notify reduce workers when task is rescheduled
- **Reduce worker failure**
  - Reset in-progress tasks to idle
  - Restart reduce task
- **Master failure**
  - Abort MapReduce task
  - Notify client

# How many Map and Reduce jobs?

---

- **M map tasks**
- **R reduce tasks**
- **N worker nodes**
- Rules of thumb
  - $M \gg N$ 
    - Pros: Improve dynamic load balancing, Speed up recovery from worker failures
    - Cons: More communication between *Master* and *Worker Nodes*, Lots of smaller files
  - $R > N$ 
    - Usually  $R < M$ , Output is spread across fewer files

# Refinements: Backup Tasks

---

- **Problem**

- Slow workers significantly lengthen the job completion time
- Slow workers due to:
  - Older processor
  - Not enough RAM
  - Other jobs on the machine
  - Bad disks
  - OS thrashing / virtual memory hell

- **Solution**

- Near the end of Map / Reduce phase
  - Spawn backup copies of tasks
  - Whichever one finishes first “wins”

- **Result**

- Shorten job completion time

# Refinement: Combiners

- **Problem**

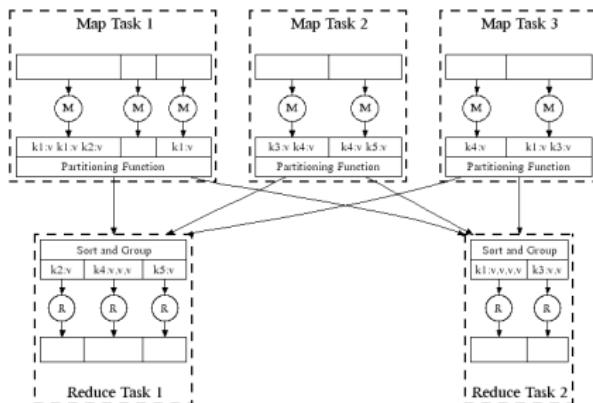
- Often a *Map* task produces many pairs for the same key  $k$   
[( $k_1$ ,  $v_1$ ), ( $k_1$ ,  $v_2$ ), ...]
- E.g., common words in the word count example
- Increase complexity of the *GroupBy* stage

- **Solution**

- Pre-aggregate values in the *Map* with a *Combine*  
[ $k_1$ , ( $v_1$ ,  $v_2$ , ...),  $k_2$ , ([...])]
- *Combine* is usually the same as the *Reduce* function
- Works only if *Reduce* function is commutative and associative

- **Result**

- Better data locality
- Less shuffling and reordering
- Less network / disk traffic



# Refinement: Partition Function

---

- **Problem**

- Sometimes users want to control how keys get partitioned
- Inputs to *Map* tasks are created by contiguous splits of input file
- MapReduce uses a default partition function **hash(key) mod R**
- Reduce needs to ensure that records with the same intermediate key end up at the same worker

- **Solution**

- Sometimes useful to override the hash function:
- E.g., **hash(hostname(URL)) mod R** ensures URLs from a host end up in the same output file

# Implementations of MapReduce

---

- There are many implementations of map reduce
- **Google**
  - Not available outside Google
- **Hadoop**
  - Open-source in Java
  - Uses HDFS for storage
  - Hadoop Wiki: Intro, Getting Started, Map/Reduce Overview
- **Amazon Elastic MapReduce (EMR)**
  - Hadoop MapReduce on Amazon EC2
  - Also runs Spark, HBase, Hive,
- **Spark**
- **Dask**