



UMD DATA605 - Big Data Systems

8.3: Apache Hadoop

- **Instructor:** Dr. GP Saggese - gsaggese@umd.edu
- **References**
 - Ghemawat et al.: *The Google File System*, 2003
 - Dean et al.: *MapReduce: Simplified Data Processing on Large Clusters*, 2004

- ***Hadoop Distributed File System***
- Hadoop MapReduce

Hadoop Distributed File System (HDFS)

- HDFS is a distributed file system
 - Designed to store large data sets reliably
 - Part of the Apache Hadoop ecosystem
 - Inspired by the Google File System (GFS)
- 1. Optimized for high-throughput access to large files
 - Suitable for batch processing
 - Not low-latency access
- 2. Designed for fault tolerance and scalability
 - Ensures fault tolerance through replication
 - Default replication factor is 3
 - Blocks are stored on different nodes and racks
 - Follows a primary-secondary architecture
 - Provides data availability even if some nodes fail
 - Replication strategy improves read performance

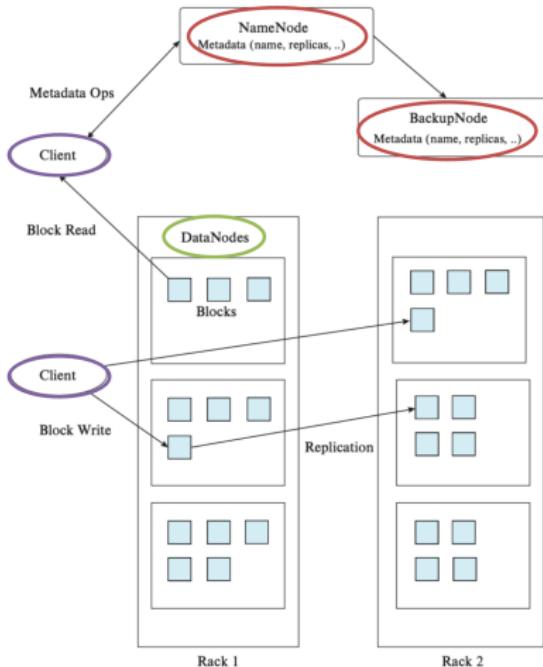
HDFS Architecture

- **NameNode**

- Store file/dir hierarchy
- Store file metadata (block location, size, permissions)

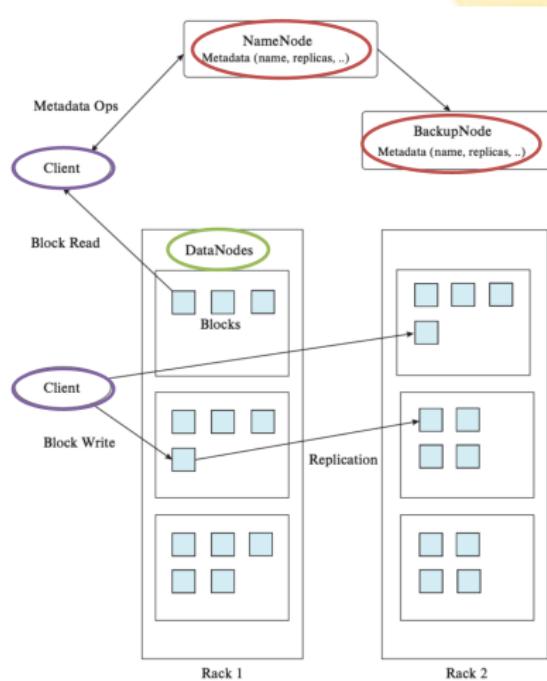
- **DataNodes**

- Store actual data blocks
- Split file into 16-64MB blocks
- Replicate chunks (2x or 3x) across multiple *DataNodes*
- Keep replicas in different racks



Hadoop Distributed File System

- Library for file access
 - Read:
 - Contact *NameNode* for *DataNode* and block pointer
 - Choose the nearest *DataNode* for each block
 - Connect to *DataNode* for data access
 - Reads blocks in parallel to improve performance
 - Data is reassembled by the client in correct order
 - Write:
 - *NameNode* creates blocks
 - Assign blocks to multiple *DataNodes*
 - Client sends data to *DataNodes*
 - *DataNodes* store



Fault Tolerance and Recovery

- *NameNode* monitors *DataNode* heartbeat signals
- On failure, blocks are re-replicated to maintain replication factor
- *NameNode* itself is a single point of failure
 - Solved with HDFS High Availability
- Data integrity ensured using checksums

HDFS vs Traditional File Systems

- Best for storing and processing large-scale logs, media, sensor data
- Commonly used in data lakes and ETL pipelines
- Optimized for write-once, read-many access pattern
- Supports very large files and directories
 - Performance degrades with many small files
- Lacks low-latency access, but provides high throughput
 - Not suitable for transactional or low-latency systems

- Hadoop Distributed File System
- *Hadoop MapReduce*

Implementations of MapReduce

- **Google**
 - Not available outside Google
- **Hadoop**
 - Open-source in Java
 - Uses HDFS for storage
 - Hadoop Wiki: Intro, Getting Started, Map/Reduce Overview
- **Amazon Elastic MapReduce (EMR)**
 - Hadoop MapReduce on Amazon EC2
 - Also runs Spark, HBase, Hive,
- **Spark**
- **Dask**

MapReduce: Hadoop

- Hadoop: open-source MapReduce implementation
- Functionalities
 - Partition input data (HDFS)
 - Input adapters
 - E.g., hBase, MongoDB, Cassandra, Amazon Dynamo
 - Schedule program execution across machines
 - Handle machine failures
 - Manage inter-machine communication
 - Perform *GroupByKey* step
 - Output adapters
 - E.g., Avro, ORC, Parquet
 - Schedule multiple *MapReduce* jobs



Data Flow

- Store input, intermediate, final outputs in HDFS
 - Operations in Hadoop move disk to disk
- Use adapters to read/partition data in chunks
- Scheduler places map tasks near physical storage of input data
 - Store intermediate results on local FS of Map and Reduce workers
- Output often serves as input for another MapReduce task

Hadoop Distributed File System (HDFS): Overview

- Designed for distributed storage of large datasets
- Built on master-slave architecture
 - NameNode: manages metadata and directory structure
 - DataNodes: store actual data blocks
- Optimized for high throughput rather than low latency
- Stores large files across multiple machines
- Writes are append-only, simplifying consistency
- Fault-tolerant using data replication
 - Default: each block is replicated 3 times
- Ideal for batch processing and big data workloads

HDFS: Data Storage and Access

- Files are split into fixed-size blocks (default: 128MB)
- Blocks distributed across DataNodes for parallelism
- NameNode maintains block-to-DataNode mapping
- Client reads data directly from DataNodes
- Ensures reliability through block replication
- If a DataNode fails, replicas serve the data
- Data locality: computation is moved to where data resides

Hadoop MapReduce: Overview

- Programming model for distributed data processing
- Processes data in parallel across a cluster
- Consists of two main functions:
 - Map: filters and sorts input data
 - Reduce: aggregates intermediate outputs
- Suited for batch jobs over large datasets
- Fault-tolerant: tasks are retried upon failure

MapReduce: Execution Phases

- Input data split into chunks processed by Mappers
- Mapper outputs key-value pairs
- Shuffle and Sort: organizes data by key
 - Intermediate keys are grouped and sent to Reducers
- Reducers aggregate values by key
- Final output written back to HDFS

Example: Word Count with MapReduce

- Input: text files containing words
- Mapper:
 - Reads lines and emits (word, 1) for each word
- Shuffle and Sort:
 - Groups by word, e.g., (word, [1,1,1])
- Reducer:
 - Sums values: emits (word, total_count)
- Output: word frequencies stored in HDFS

HDFS vs MapReduce: Complementary Roles

- HDFS: distributed storage system
 - Stores large datasets efficiently
- MapReduce: distributed compute model
 - Processes data stored in HDFS
- Together enable scalable and fault-tolerant data analysis

Benefits and Limitations

- Benefits:
 - Scalable and fault-tolerant
 - Handles petabytes of data
 - Open-source and cost-effective
- Limitations:
 - High latency, not suitable for real-time
 - Difficult for complex iterative algorithms
 - Superseded in many cases by Spark and other engines