


9.1: Apache Spark

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu
- **References:**
 - Concepts in the slides
 - Academic paper
 - “Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing”, 2012
 - Mastery
 - “Learning Spark: Lightning-Fast Data Analytics” (2nd Edition)
 - Not my favorite, but free here



Hadoop MapReduce: Shortcomings

- **Hadoop is hard to administer**


- Many layers (HDFS, Yarn, Hadoop, . . .)
- Extensive configuration

- **Hadoop is hard to use**

- Verbose API
- Limited language support (e.g., Java is native)
- MapReduce jobs read / write data on disk

- **Large but fragmented ecosystem**

- No native support for:
 - Machine learning
 - SQL, streaming
 - Interactive computing
- New systems developed on Hadoop for new workloads
- E.g., Apache Hive, Storm, Impala, Giraph, Drill



(Apache) Spark



- Open-source
 - DataBrick monetizes it (\$40B startup)
- General processing engine
 - Large set of operations beyond Map() and Reduce()
 - Combine operations in any order
 - Transformations vs Actions
 - Computation organized as a DAG
 - DAGs decomposed into parallel tasks
 - Scheduler/optimizer for parallel workers
- Supports several languages
 - Java, Scala (preferred), Python supported through bindings
- Data abstraction
 - Resilient Distributed Dataset (RDD)
 - DataFrames, Datasets built on RDDs
- Fault tolerance through RDD lineage
- In-Memory computation

Berkeley: From Research to Companies


- Amplab
- Rise lab



Berkeley AMPLab Data Analytics Stack

- So many tools that they have their own big data stack!

<https://amplab.cs.berkeley.edu/software/>




Apache Spark

- **Unified stack**
 - Different computation models in a single framework
- **Spark SQL**
 - ANSI SQL compliant
 - Work with structured relational data
- **Spark MLlib**
 - Build ML pipelines
 - Support popular ML algorithms
 - Built on top of Spark DataFrame
- **Spark Streaming**
 - Handle continually growing tables
 - Tables are treated as static table
- **GraphX**
 - Manipulate graphs
 - Perform graph-parallel computation
- **Extensibility**
 - Read from many sources
 - Write to many backends



One computation engine




General purpose applications



Resilient Distributed Dataset (RDD)

- **A Resilient Distributed Dataset (RDD)**

- Collection of data elements
- Partitioned across nodes
- Operated on in parallel
- Fault-tolerant
- In-memory / serializable



- **Applications**

- Best for applications applying the same operation to all dataset elements (vectorized)
- Less suitable for asynchronous fine-grained updates to shared state
 - E.g., updating one value in a dataframe

- **Ways to create RDDs**

- *Reference* data in external storage
 - E.g., file-system, HDFS, HBase
- *Parallelize* an existing collection in your driver program
- *Transform* RDDs into other RDDs


Transformations vs Actions

- **Transformations**

- Lazy evaluation
- Nothing computed until an Action requires it
- Build a graph of transformations

- **Actions**

- When applied to RDDs force calculations and return values
- Aka Materialize



Spark Example: Estimate Pi

```
# Estimate π (compute-intensive task).
# Pick random points in the unit square [(0,0)-(1,1)].
# See how many fall in the unit circle center=(0, 0), radius=1.
# The fraction should be π / 4.

import random
random.seed(314)

def sample(p):
    x, y = random.random(), random.random()
    in_unit_circle = 1 if x*x + y*y < 1 else 0
    return in_unit_circle

# "parallelize" method creates an RDD.
NUM_SAMPLES = int(1e6)
count = sc.parallelize(range(0, NUM_SAMPLES)) \
        .map(sample) \
        .reduce(lambda a, b: a + b)
approx_pi = 4.0 * count / NUM_SAMPLES
print("pi is roughly %f" % approx_pi)
```

executed in 386ms, finished 04:27:53 2022-11-23

pi is roughly 3.141400

SCIENCE
ACADEMY



Spark: Architecture

- **Architecture**

- Who does what
- Responsibilities of each piece

- **Spark Application**

- Code describing computation
- E.g., Python code calling Spark

- **Spark Driver**

- Instantiate *SparkSession*
- Communicate with *Cluster Manager* for resources
- Transform operations into DAG computations
- Distribute task execution across *Executors*

- **Spark Session**


- Interface to Spark system

- **Cluster Manager**

- Manage and allocate resources
- Support Hadoop, YARN, Mesos, Kubernetes


- **Spark Executor**

- Run worker node to execute tasks
- Typically one executor per node



Spark: Computation Model

- **Architecture** = who does what
- **Computational model** = how are things done




- **Spark Driver**
 - Converts Spark application into one or more Spark *Jobs*
 - Describes computation with *Transformations* and triggers with *Actions*
- **Spark Job**
 - Parallel computation in response to a Spark *Action*
 - Each *Job* is a DAG with one or more dependent *Stages*
- **Spark Stage**
 - Smaller operation within a *Job*
 - Stages can run serially or in parallel
- **Spark Task**
 - Each Stage has multiple *Tasks*

Deployment Modes

- Spark can run on several different configurations
 - **Local**
 - E.g., run on your laptop
 - Driver, Cluster Manager, Executors all run in a single JVM on the same node
 - **Standalone**
 - Driver, Cluster Manager, Executors run in different JVMs on different nodes
 - **YARN or Kubernetes**
 - Driver, Cluster Manager, Executors run on different pods (i.e., containers)


Distributed Data and Partitions

- **Data is distributed** as partitions across different physical nodes
 - Store each partition in memory
 - Enable efficient parallelism
- **Spark Executors** process data “close” to them
 - Minimize network bandwidth
 - Ensure data locality
 - Similar approach to Hadoop



Parallelized Collections


- Parallelized collections created by calling *SparkContext parallelize()* on an existing collection



- Data spread across nodes
- Number of *partitions* to cut dataset into
 - Spark runs one *Task* per partition
 - Aim for 2-4 partitions per CPU
 - Spark sets partitions automatically based on your cluster
 - Set manually by passing as a second parameter to `parallelize()`

Transformations vs Actions

- Transform a Spark RDD into a new RDD without modifying the input data
 - Immutability like in functional programming
 - E.g., `select()`, `filter()`, `join()`, `orderBy()`
- Transformations are evaluated lazily
 - Inspect computation and decide how to optimize it
 - E.g., joining, pipeline operations, breaking into stages
- Results are recorded as “lineage”
 - A sequence of stages that can be rearranged, optimized without changing results
- **Actions**
- An action triggers the evaluation of a computation
 - E.g., `show()`, `take()`, `count()`, `collect()`, `save()`



Spark Example: MapReduce in 1 or 4 Line

```
!more data.txt
```

executed in 1.77s, finished 04:37:35 2022-11-23

One a penny, two a penny, hot cross buns

```
lines = sc.textFile("data.txt").flatMap(lambda line: line.split(" "))
pairs = lines.map(lambda s: (s, 1))
counts = pairs.reduceByKey(lambda a, b: a + b)
result = counts.collect()
print(result)
```

executed in 428ms, finished 04:36:24 2022-11-23

[('One', 1), ('two', 1), ('hot', 1), ('cross', 1), ('a', 2), ('penny,', 2), ('buns', 1)]

4 lines

MapReduce in

```
result = sc.textFile("data.txt").flatMap(lambda line: line.split(" ")).map(
    lambda s: (s, 1)).reduceByKey(lambda a, b: a + b).collect()
print(result)
```

executed in 591ms, finished 05:06:00 2022-11-23

[('One', 1), ('two', 1), ('hot', 1), ('cross', 1), ('a', 2), ('penny,', 2), ('buns', 1)]

1 line (*show-off version*)

MapReduce in

Same Code in Java Hadoop

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }
}
```

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
                      Context context
                      throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```



Spark Example: Logistic Regression in MapReduce

- Logistic Regression (from [here](#))

```
# Load points
points = spark.textFile(...).map(parsePoint)

# Initial separating plane
w = numpy.random.ranf(size=D)

# Until convergence
for i in range(ITERATIONS):
    # Parallel loop over the samples i=1...m
    gradient = points.map(
        lambda p: (1 / (1 + exp(-p.y*(w.dot([
            .reduce(lambda a, b: a + b)
    w -= alpha * gradient

print("Final separating plane: %s" % w)
```

Repeat {

$$\theta_j := \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2$$

Repeat until convergence {

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

Spark Transformations: 1 / 3

- `map(func)`
 - Return new RDD passing each element through `func()`
- `flatMap(func)`
 - Map each input item to 0 or more output items
 - `func()` returns a sequence
- `filter(func)`
 - Return new RDD selecting elements where `func()` returns true
- `union(otherDataset)`
 - Return new RDD with union of elements in source dataset and argument
- `intersection(otherDataset)`
 - Return new RDD with intersection of elements in source dataset and argument

<https://spark.apache.org/docs/latest/rdd-programming-guide.html>

Spark Transformations: 2 / 3

- `distinct([numTasks])`
 - Return new RDD with distinct elements of source dataset
- `join(otherDataset, [numTasks])`
 - On RDDs (K, V) and (K, W) , return dataset of $(K, (V, W))$ pairs for each key
 - Support outer joins: `leftOuterJoin`, `rightOuterJoin`, `fullOuterJoin`
- `cogroup(otherDataset, [numPartitions])`
 - Aka `groupWith()`
 - Like join but return dataset of $(K, (Iterable<V>, Iterable<W>))$ tuples

Spark Transformations: 3 / 3

- `groupByKey([numPartitions])`
 - On RDD of (K, V) pairs, returns (K, Iterable<V>) pairs
 - For aggregation (e.g., sum, average), use `reduceByKey` for better performance
 - Process data in place instead of iterators
 - Output parallelism depends on parent RDD partitions
 - Use `numPartitions` to set tasks
- `reduceByKey(func, [numPartitions])`
 - On RDD of (K, V) pairs, returns (K, f(V_1, ..., V_n)) pairs with values aggregated by `func()`
 - `func(): (V, V) → V`
 - Shuffle + Reduce from MapReduce
 - Configure reduce tasks with `numPartitions`
- `sortByKey([ascending], [numPartitions])`
 - Returns (K, V) pairs sorted by keys in ascending or descending order


Spark Actions

- `reduce(func)`
 - Aggregate dataset elements using `func()`
 - `func()` takes two arguments, returns one
 - `func()` must be commutative and associative for parallel computation
- `collect()`
 - Return dataset elements as an array
 - Useful after operations returning a small data subset (e.g., `filter()`)
- `count()`
 - Return number of elements in the dataset
- `take(n)`
 - Return array with first `n` dataset elements
 - `.collect()[:n]` differs from `.take(n)`

<https://spark.apache.org/docs/latest/rdd-programming-guide.html>

Spark: Fault-tolerance


- Spark uses *immutability* and *lineage* for fault tolerance
- In case of failure:
 - Reproduce RDD by replaying lineage
 - No need for checkpoints
 - Keep data in memory to boost performance
- Fault-tolerance is free!



Spark: RDD Persistence

- User explicitly persists (aka cache) an RDD

- Call `persist()`, `unpersist()` on RDD
- Cache if RDD is expensive to compute
 - E.g., filtering large data
- When you persist an RDD, each node:
 - Stores (in memory or disk) partitions of the RDD
 - Reuses cached partitions on derived datasets



- Cache

- Makes future actions faster (often >10x)
- Managed by Spark with LRU policy + garbage collector

- User can choose storage level

- MEMORY_ONLY (default)
- DISK_ONLY (e.g., Python Pickle)
- MEMORY_AND_DISK
 - If RDD doesn't fit in memory, store on disk
- MEMORY_AND_DISK_2


- Same as above, replicate each partition on two nodes

Caching on disk can be more expensive than not caching




Spark: RDD Persistence and Fault-tolerance

- Spark handles persistence and fault-tolerance similarly
- **Caching/Persistence**
 - Cache RDD (in memory or on disk) instead of recomputing
- **Fault-tolerance**
 - If any partition of an RDD is lost
 - Automatically recompute RDD using transformations that generated it
 - Based on immutability and lineage
- **Caching is fault-tolerant!**



Spark Shuffle

- E.g., **reduceByKey()**
 - *Definition:* Combine values $[v_1, \dots, v_n]$ for key k into (k, v) where $v = \text{reduce}(v_1, \dots, v_n)$
 - *Problem:* Values for a key must be on the same partition/machine
 - *Solution:* Shuffle data across machines
- Certain Spark operations trigger a data shuffle
 - E.g., `reduceByKey()`, `groupByKey()`, `join`, `repartition`, `transpose`
- **Data shuffle** = Re-distribute data across partitions/machines
- **Data shuffle is expensive** due to:
 - Data serialization (pickle)
 - Disk I/O (save to disk)
 - Network I/O (copy across Executors)
 - Deserialization and memory allocation
- **Spark schedules general task graphs**
 - Automatic function pipelining
 - Data locality aware



Broadcast Variables

- **Problem**

- Ship common variables to nodes with code
- Broadcasting involves serialization, network transfer, de-serialization
- Sending large, constant data repeatedly is costly

- **Solution**

- Cache read-only variables on each node, avoid task copies

```
# `var` is large variable.
```

```
var = list(range(1, int(1e6)))
```

```
# Create a broadcast variable.
```

```
broadcast_var = sc.broadcast(var)
```

```
# Do not modify `var`, but use `broadcast_var.value` instead of var
```

Accumulators

- **Accumulator** = variable “added to” through associative, commutative operations

- Efficient in parallel execution (e.g., MapReduce)

- Spark supports Accumulators with numerical types (e.g., integers)
 - Define Accumulators for different types

```
>>> accum = sc.accumulator(0)
```

```
>>> accum
```

```
Accumulator<id=0, value=0>
```

```
>>> sc.parallelize([1, 2, 3, 4]).foreach(lambda x: accum.add(x))
```

```
>>> accum.value
```

```
10
```

- Each node computes value to add to Accumulator

- Usual semantic:

- Accumulators use logic of transformations (lazy evaluation) and actions

```
accum = sc.accumulator(0)
```

```
def g(x):
```

```
    accum.add(x)
```

```
    return f(x)
```

```
data.map(g)
```

Here, accum is still 0 because no actions have caused the

Gray Sort Competition

	Hadoop MR Record	Spark Record (2014)
Data Size	102.5 TB	100 TB
Elapsed Time	72 mins	23 mins
# Nodes	2100	206
# Cores	50400 physical	6592 virtualized
Cluster disk throughput	3150 GB/s	618 GB/s
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network
Sort rate	1.42 TB/min	4.27 TB/min
Sort rate/node	0.67 GB/min	20.7 GB/min

- From [here](#)
- Sort benchmark, Daytona Gray: sort of 100 TB of data (1 trillion records)
 - Spark-based System 3x faster with 1/10 number of nodes

Spark vs Hadoop MapReduce

- **Performance:** Spark normally faster but with caveats
 - Spark can process data in-memory
 - Spark generally outperforms MapReduce, but it often needs lots of memory to do well
 - Hadoop MapReduce persists back to the disk after a map or reduce action
- **Ease of use:** Spark is easier to program
- **Data processing:** Spark more general

“Spark vs. Hadoop MapReduce”, Saggi Neumann, 2014

<https://www.xplenty.com/blog/2014/11/apache-spark-vs-hadoop-mapreduce/>